

## Week 3

### Policy

A **policy**  $\pi$  is a mapping from states to probabilities of selecting each possible action:

$$\pi(a | s) = \Pr(A_t = a | S_t = s)$$

Type	Description
Deterministic	$\pi(s) = a$ — always selects the same action in state $s$
Stochastic	$\pi(a   s)$ — probability distribution over actions

### Value Functions

#### State-Value Function

$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s]$$

- Expected return starting from state  $s$  and following policy  $\pi$
- Measures how good state  $s$  is **under policy**  $\pi$

#### Action-Value Function

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$$

- Expected return starting from state  $s$ , taking action  $a$ , then following policy  $\pi$
- Measures how good taking action  $a$  in state  $s$  is **under policy**  $\pi$

#### Relationship Between $v_\pi$ and $q_\pi$

From	To	Formula
$q_\pi$	$v_\pi$	$v_\pi(s) = \sum_a \pi(a   s) q_\pi(s, a)$
$v_\pi$	$q_\pi$	$q_\pi(s, a) = \sum_{s', r} p(s', r   s, a) [r + \gamma v_\pi(s')]$

Here  $r$  denotes possible values of the random reward  $R_{t+1}$ .

### Value Functions Satisfy Recursive Relationships

Value functions can be expressed recursively — the value of the current state depends on the values of successor states. This recursive structure arises from the definition of return:

$$G_t = R_{t+1} + \gamma G_{t+1}$$

Since the return  $G_t$  is defined as the immediate reward plus the discounted future return, we can write:

$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s] = \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} | S_t = s]$$

### Markov Property Enables Recursive Closure

The key step from  $G_{t+1}$  to  $v_\pi(S_{t+1})$  relies on the **Markov property**: the distribution of  $G_{t+1}$  only depends on the next state  $S_{t+1}$ , not on earlier history.

$$\mathbb{E}_\pi[G_{t+1} \mid S_t = s, A_t = a, S_{t+1} = s'] = \mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s'] = v_\pi(s')$$

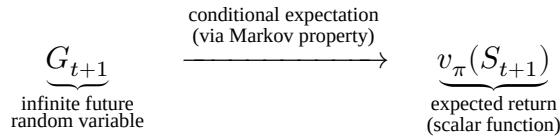
This allows us to substitute  $G_{t+1}$  with  $v_\pi(S_{t+1})$ , yielding the **recursive form**:

$$v_\pi(s) = \mathbb{E}_\pi[R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s]$$


---

### The Key Compression: Infinite Future $\rightarrow$ Single Scalar

This is the **mathematical foundation of reinforcement learning**:



More precisely:  $\mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s'] = v_\pi(s')$

Aspect	$G_{t+1}$	$v_\pi(S_{t+1})$
Nature	Random variable (sum of infinite rewards)	Scalar function = <b>expected return</b>
Depends on	Entire future trajectory	Only current state $S_{t+1}$
Computational	Impossible to compute directly	Can be estimated recursively
Role	Definition of return	<b>Compression</b> via conditional expectation

**Why this compression is profound:** Without the Markov property,  $\mathbb{E}_\pi[G_{t+1} \mid S_t, A_t, S_{t+1}, R_{t+1}]$  would depend on the entire history, making recursive computation impossible. The Markov property enables us to **discard all past conditions** and keep only  $S_{t+1}$ .

This compression is what makes the following possible:

- **Dynamic Programming:** Recursively solve for value functions
- **Temporal-Difference Learning:** Learn from incomplete episodes
- **Generalization:** Approximate value functions with neural networks

The value of a state can be decomposed into the **immediate reward** received after leaving that state, plus the **discounted value** of the successor state.

Component	Meaning
$R_{t+1}$	Immediate reward after taking action in state $s$
$\gamma G_{t+1}$	Discounted future return from successor state $s'$
$\gamma v_\pi(S_{t+1})$	Discounted value of successor state (by Markov property)

This recursive property is fundamental because:

- Enables bootstrapping:** We can estimate the value of a state using estimates of successor states, without waiting for the episode to end
- Forms the basis of Bellman equations:** The recursive relationship is formalized into equations that relate values across states
- Powers DP and TD methods:** Dynamic programming and temporal-difference learning exploit this structure for efficient computation

## Bellman Equations

**Bellman Equation for  $v_\pi$**

$$v_\pi(s) = \sum_a \pi(a | s) \sum_{s',r} p(s',r | s,a) [r + \gamma v_\pi(s')]$$

The value of a state equals the expected immediate reward plus the discounted value of the next state.

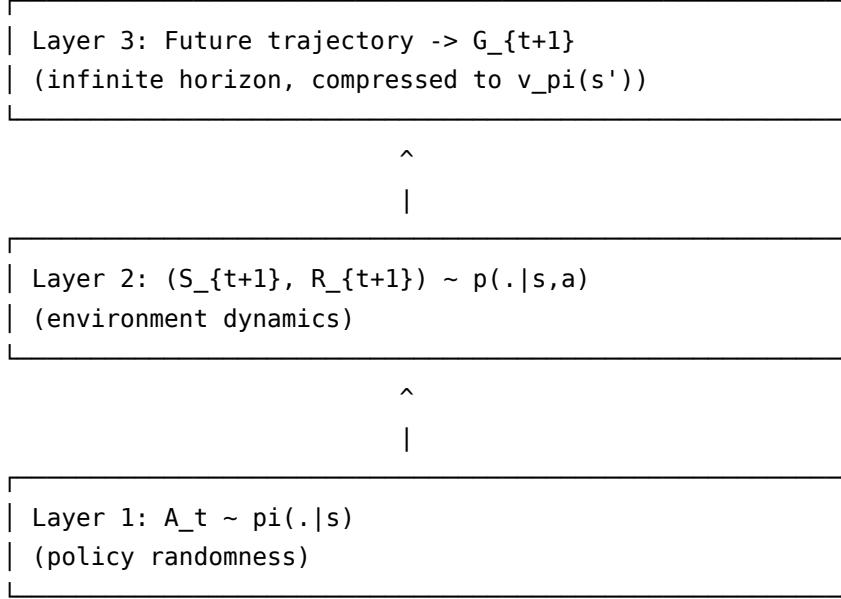
## Derivation:

Starting point:  $v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s]$  and  $G_t = R_{t+1} + \gamma G_{t+1}$

---

### The Three-Layer Structure of Randomness

To compute  $\mathbb{E}_\pi[G_t | S_t = s]$ , we must average over **three layers of randomness**:



Each application of the **law of total expectation** "peels off" one layer.

Layer	Random variable	Distribution	What we average over
1	Action $A_t$	$\pi(a   s)$	Policy's choice
2	Next state/reward $(S_{t+1}, R_{t+1})$	$p(s',r   s,a)$	Environment's response
3	Future return $G_{t+1}$	Trajectory distribution given $S_{t+1} = s'$	<b>Compressed</b> by Markov property: $\mathbb{E}_\pi[G_{t+1}   S_{t+1} = s']$

The **Markov property** is the key that enables Layer 3 compression: the infinite future  $G_{t+1}$  is reduced to a single scalar  $v_\pi(S_{t+1})$  through conditional expectation.

---

**Step 1 — Peel Layer 1:** Condition on action  $A_t$  (law of total expectation):

$$\begin{aligned} v_\pi(s) &= \sum_a \underbrace{\Pr(A_t = a \mid S_t = s)}_{\pi(a|s)} \cdot \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a] \\ &= \sum_a \pi(a \mid s) \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid s, a] \end{aligned}$$

(Shorthand:  $\mathbb{E}_\pi[\cdot \mid s, a] \equiv \mathbb{E}_\pi[\cdot \mid S_t = s, A_t = a]$ )

---

**Step 2 — Peel Layer 2:** Condition on  $(S_{t+1}, R_{t+1})$ :

$$\begin{aligned} \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid s, a] &= \sum_{s',r} \underbrace{\Pr(S_{t+1} = s', R_{t+1} = r \mid s, a)}_{p(s',r \mid s, a)} \\ &\quad \times \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a, S_{t+1} = s', R_{t+1} = r] \\ &= \sum_{s',r} p(s',r \mid s, a) [r + \gamma \mathbb{E}_\pi[G_{t+1} \mid S_t = s, A_t = a, S_{t+1} = s', R_{t+1} = r]] \end{aligned}$$

(Since  $R_{t+1} = r$  is fixed, it comes out of the expectation as  $r$ )

---

**Step 3 — Layer 3 Compression:** Apply Markov property

By the Markov property,  $G_{t+1}$  depends only on  $S_{t+1} = s'$ , **not on earlier history** ( $S_t = s, A_t = a, R_{t+1} = r$ ):

$$\mathbb{E}_\pi[G_{t+1} \mid S_t = s, A_t = a, S_{t+1} = s', R_{t+1} = r] = \mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s']$$

**This is the key compression:** We can **delete all past conditions** and keep only  $S_{t+1}$ .

By the **definition of state-value function** (at time  $t + 1$ , by time-homogeneity):

$$\mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s'] = v_\pi(s')$$

**Infinite future**  $G_{t+1}$  has been **compressed to a single scalar**  $v_\pi(s')$ .

---

**Step 4 — Combine all steps:**

$$v_\pi(s) = \sum_a \pi(a \mid s) \sum_{s',r} p(s',r \mid s, a) [r + \gamma v_\pi(s')]$$

---

**Intuition:** The value of state  $s$  averages over:

1. **Actions** (weighted by policy  $\pi$ )
2. **Outcomes**  $(s', r)$  (weighted by dynamics  $p$ )
3. **Immediate reward  $r$  + discounted future value**  $\gamma v_\pi(s')$

**Bellman Equation for  $q_\pi$**

$$q_\pi(s, a) = \sum_{s', r} p(s', r | s, a) \left[ r + \gamma \sum_{a'} \pi(a' | s') q_\pi(s', a') \right]$$

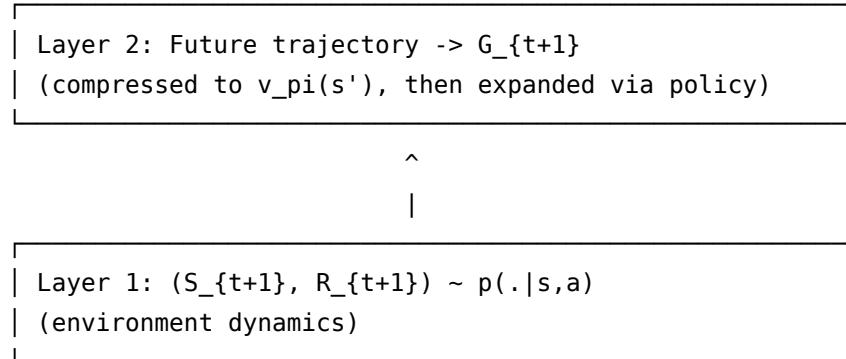
**Derivation:**

Starting point:  $q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$  and  $G_t = R_{t+1} + \gamma G_{t+1}$

**Note:** For  $q_\pi$ , the action  $A_t = a$  is **already fixed**, so we have only **two layers** of randomness (not three).

---

**The Two-Layer Structure for  $q_\pi$**



Layer	Random variable	Distribution
1	Next state/reward $(S_{t+1}, R_{t+1})$	$p(s', r   s, a)$
2	Future return $G_{t+1}$	Compressed to $v_\pi(s')$ , then expanded: $v_\pi(s') = \sum_{a'} \pi(a'   s') q_\pi(s', a')$

---

**Step 1 — Peel Layer 1:** Condition on  $(S_{t+1}, R_{t+1})$ :

$$\begin{aligned} q_\pi(s, a) &= \sum_{s', r} \underbrace{\Pr(S_{t+1} = s', R_{t+1} = r | s, a)}_{p(s', r | s, a)} \cdot \mathbb{E}_\pi[G_t | s, a, s', r] \\ &= \sum_{s', r} p(s', r | s, a) \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} | S_t = s, A_t = a, S_{t+1} = s', R_{t+1} = r] \\ &= \sum_{s', r} p(s', r | s, a) [r + \gamma \mathbb{E}_\pi[G_{t+1} | S_t = s, A_t = a, S_{t+1} = s', R_{t+1} = r]] \end{aligned}$$

(Since  $R_{t+1} = r$  is fixed, it comes out of the expectation as  $r$ )

---

**Step 2 — Layer 2 Compression:** Apply Markov property

By the Markov property,  $G_{t+1}$  depends only on  $S_{t+1} = s'$ , **not on earlier history**:

$$\mathbb{E}_\pi[G_{t+1} \mid S_t = s, A_t = a, S_{t+1} = s', R_{t+1} = r] = \mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s']$$

By the **definition of state-value function** (at time  $t + 1$ , by time-homogeneity):

$$\mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s'] = v_\pi(s')$$

Infinite future  $G_{t+1}$  compressed to scalar  $v_\pi(s')$ .

So we have the intermediate form:

$$q_\pi(s, a) = \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_\pi(s')]$$


---

**Step 3 — Expand  $v_\pi$  back to  $q_\pi$ :**

Apply law of total expectation to  $v_\pi(s') = \mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s']$  by conditioning on the **next action**  $A_{t+1}$ :

$$\begin{aligned} v_\pi(s') &= \sum_{a'} \underbrace{\Pr(A_{t+1} = a' \mid S_{t+1} = s')}_{\pi(a' \mid s')} \cdot \mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s', A_{t+1} = a'] \\ &= \sum_{a'} \pi(a' \mid s') \cdot q_\pi(s', a') \end{aligned}$$

(By time-homogeneity,  $\mathbb{E}_\pi[G_{t+1} \mid S_{t+1} = s', A_{t+1} = a'] = q_\pi(s', a')$  — the action-value function definition applies at any time step)

**Compression followed by expansion:**  $G_{t+1} \rightarrow v_\pi(s') \rightarrow \sum \pi(\cdot) q_\pi(\cdot)$

---

**Step 4 — Substitute  $v_\pi(s')$  back:**

$$q_\pi(s, a) = \sum_{s', r} p(s', r \mid s, a) \left[ r + \gamma \sum_{a'} \pi(a' \mid s') q_\pi(s', a') \right]$$

---

**Intuition:** The value of action  $a$  in state  $s$  averages over:

1. **Outcomes**  $(s', r)$  (weighted by dynamics  $p$ )
2. **Immediate reward**  $r + \text{discounted value of next state}$   $\gamma v_\pi(s')$
3. Where  $v_\pi(s')$  itself averages over **next actions** (weighted by policy  $\pi$ )

## Backup Diagrams

Diagram	Description
State-value backup	Shows how $v_\pi(s)$ depends on $v_\pi(s')$ for successor states
Action-value backup	Shows how $q_\pi(s, a)$ depends on $q_\pi(s', a')$ for successor state-action pairs

- White circle: state node
- Black dot: action node
- Arcs from state nodes represent policy  $\pi(a | s)$
- Arcs from action nodes represent dynamics  $p(s', r | s, a)$

## Optimal Value Functions

### Optimal State-Value Function

$$v_*(s) = \max_{\pi} v_{\pi}(s), \quad \forall s \in \mathcal{S}$$

### Optimal Action-Value Function

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a), \quad \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$$

### Relationship

$$v_*(s) = \max_a q_*(s, a)$$

$$q_*(s, a) = \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]$$

## Bellman Optimality Equations

### For $v_*$

$$v_*(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]$$

### For $q_*$

$$q_*(s, a) = \sum_{s', r} p(s', r | s, a) \left[ r + \gamma \max_{a'} q_*(s', a') \right]$$

## Optimal Policy

A policy  $\pi$  is optimal if  $v_\pi(s) \geq v_{\pi'}(s)$  for all states  $s$  and all policies  $\pi'$ .

Property	Description
Existence	At least one optimal policy always exists
Shared value	All optimal policies share the same $v_*$ and $q_*$
Greedy extraction	Given $q_*$ , optimal policy is $\pi_*(s) = \arg \max_a q_*(s, a)$